

Introduction to Deep Learning

Yixin Lin, Serge Assaad, Rohith Kuditipudi

Duke University

yixin.lin@duke.edu, serge.assaad@duke.edu, rohith.kuditipudi@duke.edu

March 25, 2017

Overview


- 1 Review of basic principles of machine learning
- 2 Introduction to neural networks
- 3 Tutorial on implementing deep learning algorithms

Review of basic principles of machine learning

- Three components to any ML problem: the **task**, the **performance measure** and the **data**
- Essential definitions
 - **Features**
 - **Model**
 - **Parameters**



Example: Linear Regression

- Can we **predict** y given x , using the model $\hat{y} = mx + b$?




Introduction to Neural Networks

- Neurons are the building blocks of neural networks
- Each neuron is a **function**: $y = f(\mathbf{w}^T \mathbf{x} + \mathbf{b})$



Neural Networks are Layers of Neurons



What are Neural Networks Good For?

airplane




automobile




bird



cat




deer



dog



frog



horse





ship



truck



Training Neural Networks: Gradient Descent



- Time to **build something!**

Thanks!

Resources and references: yixinlin.net/intro-ml